

Demand Forecasting for a Fast-Food Restaurant Chain

Logistics and Supply Chain Analytics - Individual Project

Konstantinos Paganopoulos

Solutions

We first load the necessary libraries.

We have a dataset, which includes daily sales for lettuce at a store in New York from a fast-food restaurant chain from 5 March 2015 to 15 June 2015. Each observation includes two values: day pair, and sales in that particular day.

Then, we load and split the data set into train and test set.

```
# read csv file
data <- read.csv(file = "NewYork1_final.csv", header = TRUE, stringsAsFactors = FALSE)

# convert column date of data set to type date
data$date <- as.Date(data$date)

# convert sales into a time series object
lettuce <- ts(data[, 2], frequency = 7, start = c(10, 1)) # 10th week 1st day


# split data set into train and test set
lettuce_train <- subset(lettuce, end = 89)
lettuce_test <- subset(lettuce, start = 90) # last 14 lines-days
```

ARIMA


We visually inspect the time series.

```
autoplot(lettuce_train, xlab = "Weeks", ylab = "Lettuce") + theme_minimal() +
ggtitle("New York 1 Store (12631) - Time series plot")
```

New York 1 Store (12631) – Time series plot



```
ggtsdisplay(lettuce_train, xlab = "weeks", theme = theme_minimal())
```



Due to the trend in time series, it is non-stationary. We can get rid of trend by taking first-order difference. We plot the time series after the difference, and observe that there is no trend and appears to be stationary. We run ADF, PP and KPSS tests to formally test the stationarity of time series after the first-order difference, and all suggest that the time series is stationary.

```
# stationary test
adf.test(lettuce_train)
```

```
## Warning in adf.test(lettuce_train): p-value smaller than printed p-value
```

```
##
## Augmented Dickey-Fuller Test
##
## data: lettuce_train
## Dickey-Fuller = -4.8181, Lag order = 4, p-value = 0.01
## alternative hypothesis: stationary
```

```
pp.test(lettuce_train)
```

```
## Warning in pp.test(lettuce_train): p-value smaller than printed p-value
```

```
##
## Phillips-Perron Unit Root Test
##
```

```
## data: lettuce_train
## Dickey-Fuller Z(alpha) = -75.753, Truncation lag parameter = 3, p-value
## = 0.01
## alternative hypothesis: stationary
```

```
kpss.test(lettuce_train)
```

```
## Warning in kpss.test(lettuce_train): p-value smaller than printed p-value
```

```
##
## KPSS Test for Level Stationarity
##
## data: lettuce_train
## KPSS Level = 0.96112, Truncation lag parameter = 3, p-value = 0.01
```

The two automatic functions, `ndiffs()` and `nsdiffs()` tell us how many first-order differences, and how many seasonal differences, respectively, we need to take to make the time series stationary. We use those functions below:

```
ndiffs(lettuce_train)
```

```
## [1] 1
```

```
# seasonal stationarity
nsdiffs(lettuce_train)
```

```
## [1] 0
```

We need to differentiate one time.

```
### stationarize time series
# take first order difference
lettuce_train.diff1 <- diff(lettuce_train, differences = 1)
```

Check again the tests for stationarity:

```
# stationary test
adf.test(lettuce_train.diff1)
```

```
## Warning in adf.test(lettuce_train.diff1): p-value smaller than printed p-value
```

```
##
## Augmented Dickey-Fuller Test
##
## data: lettuce_train.diff1
## Dickey-Fuller = -7.178, Lag order = 4, p-value = 0.01
## alternative hypothesis: stationary
```

```
pp.test(lettuce_train.diff1)
```

```
## Warning in pp.test(lettuce_train.diff1): p-value smaller than printed p-value
```

```
##  
## Phillips-Perron Unit Root Test  
##  
## data: lettuce_train.diff1  
## Dickey-Fuller Z(alpha) = -112.24, Truncation lag parameter = 3, p-value  
## = 0.01  
## alternative hypothesis: stationary
```

```
kpss.test(lettuce_train.diff1)
```

```
## Warning in kpss.test(lettuce_train.diff1): p-value greater than printed p-value
```

```
##  
## KPSS Test for Level Stationarity  
##  
## data: lettuce_train.diff1  
## KPSS Level = 0.046459, Truncation lag parameter = 3, p-value = 0.1
```

Check again the two automatic functions for stationarity:

```
ndiffs(lettuce_train.diff1)
```

```
## [1] 0
```


```
nsdiffs(lettuce_train.diff1)
```

```
## [1] 0
```

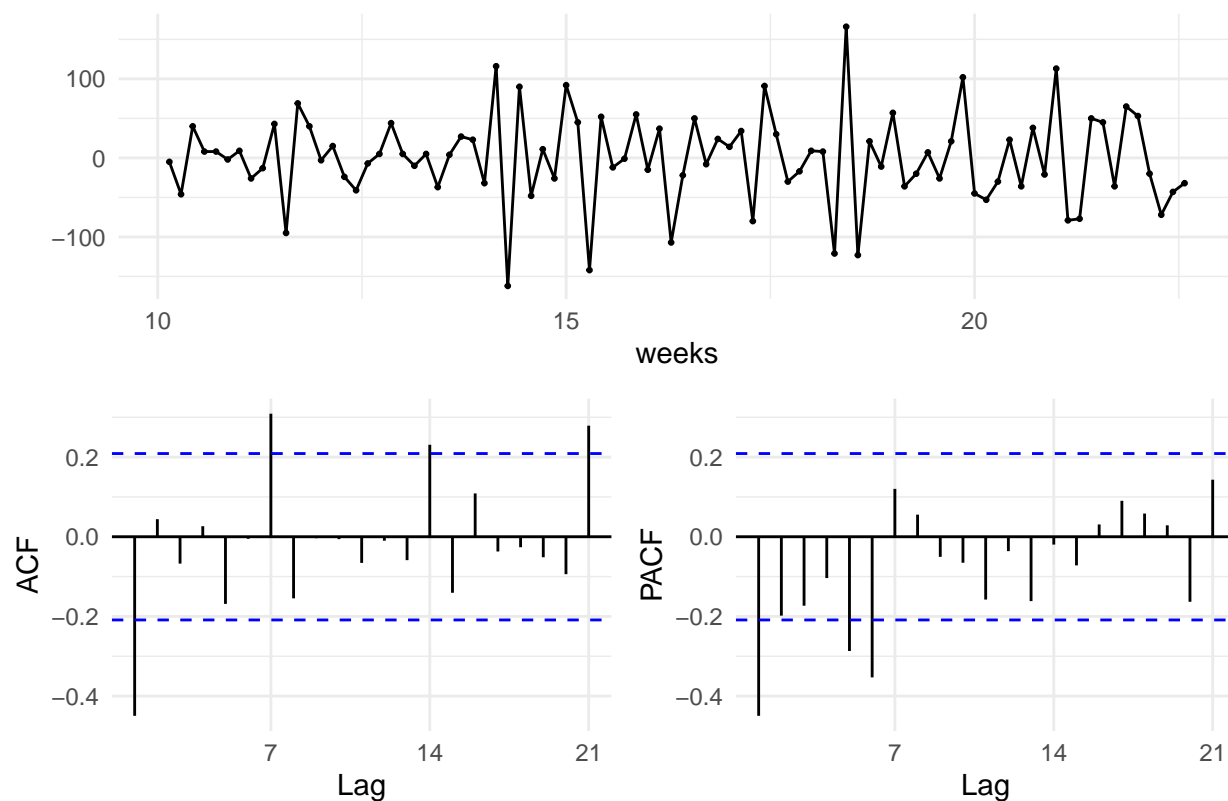
We now visually inspect the differentiated time series.

```
autoplot(lettuce_train.diff1, xlab = "Weeks", ylab = "Lettuce") + theme_minimal() +  
ggtitle("New York 1 Store (12631) - Time series plot (Differentiated)")
```

New York 1 Store (12631) – Time series plot (Differentiated)



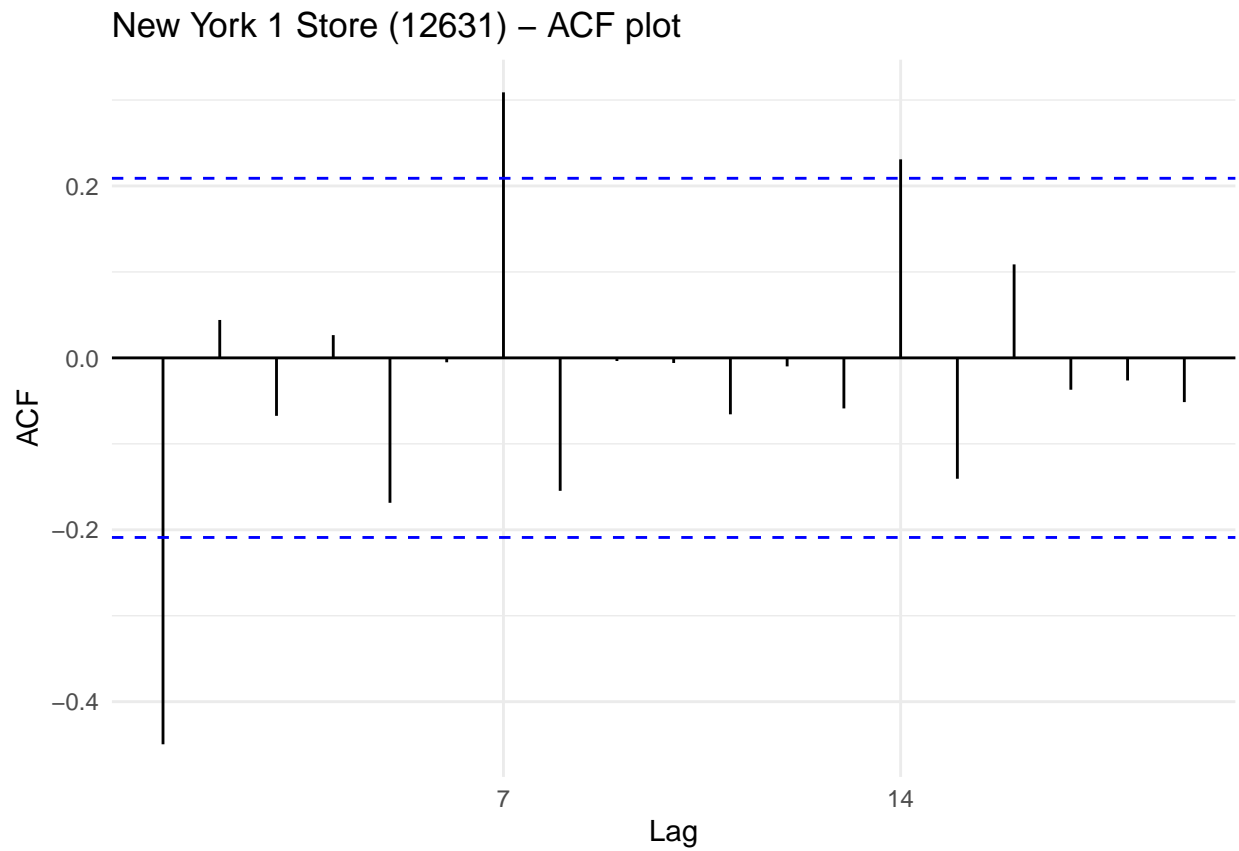
```
ggtsdisplay(lettuce_train.diff1, xlab = "weeks", theme = theme_minimal())
```




Looks stationary.

Once we have a stationary time series, the next step is to determine the optimal orders of MA and AR components. We first plot the ACF and PACF of the time series.

```
# acf plot
ggAcf(lettuce_train.diff1) + theme_minimal() + ggtitle("New York 1 Store (12631) - ACF plot")
```



```
# pacf plot  
ggPacf(lettuce_train.diff1) + theme_minimal() + ggtitle("New York 1 Store (12631) - PACF plot")
```

Next we use `auto.arima()` to search for the best ARIMA models.

The default procedure uses some approximations to speed up the search. These approximations can be avoided with the argument `approximation = FALSE`. It is possible that the minimum AIC model will not be found due to these approximations, or because of the stepwise procedure. A much larger set of models will be searched if the argument `stepwise = FALSE` is used. We also use `d = 1` and `D = 0` since we had first-differencing but no seasonal-differencing.

```
auto.arima(lettuce_train, trace = TRUE, ic = 'aic', approximation = FALSE, stepwise = FALSE, d=1, D=0)
```

```
##
## ARIMA(0,1,0) : 962.3509
## ARIMA(0,1,0) with drift : 964.3481
## ARIMA(0,1,0)(0,0,1)[7] : 957.5856
## ARIMA(0,1,0)(0,0,1)[7] with drift : 959.5792
## ARIMA(0,1,0)(0,0,2)[7] : 958.2074
## ARIMA(0,1,0)(0,0,2)[7] with drift : 960.2018
## ARIMA(0,1,0)(1,0,0)[7] : 955.5445
## ARIMA(0,1,0)(1,0,0)[7] with drift : 957.5381
## ARIMA(0,1,0)(1,0,1)[7] : Inf
## ARIMA(0,1,0)(1,0,1)[7] with drift : Inf
## ARIMA(0,1,0)(1,0,2)[7] : Inf
## ARIMA(0,1,0)(1,0,2)[7] with drift : Inf
## ARIMA(0,1,0)(2,0,0)[7] : 955.1941
## ARIMA(0,1,0)(2,0,0)[7] with drift : 957.1903
## ARIMA(0,1,0)(2,0,1)[7] : Inf
```

```

## ARIMA(0,1,0)(2,0,1)[7] with drift : Inf
## ARIMA(0,1,0)(2,0,2)[7] : Inf
## ARIMA(0,1,0)(2,0,2)[7] with drift : Inf
## ARIMA(0,1,1) : 923.6733
## ARIMA(0,1,1) with drift : 923.4671
## ARIMA(0,1,1)(0,0,1)[7] : 919.049
## ARIMA(0,1,1)(0,0,1)[7] with drift : Inf
## ARIMA(0,1,1)(0,0,2)[7] : 917.2619
## ARIMA(0,1,1)(0,0,2)[7] with drift : Inf
## ARIMA(0,1,1)(1,0,0)[7] : 915.5441
## ARIMA(0,1,1)(1,0,0)[7] with drift : Inf
## ARIMA(0,1,1)(1,0,1)[7] : Inf
## ARIMA(0,1,1)(1,0,1)[7] with drift : Inf
## ARIMA(0,1,1)(1,0,2)[7] : Inf
## ARIMA(0,1,1)(1,0,2)[7] with drift : Inf
## ARIMA(0,1,1)(2,0,0)[7] : 912.3795
## ARIMA(0,1,1)(2,0,0)[7] with drift : Inf
## ARIMA(0,1,1)(2,0,1)[7] : Inf
## ARIMA(0,1,1)(2,0,1)[7] with drift : Inf
## ARIMA(0,1,1)(2,0,2)[7] : Inf
## ARIMA(0,1,1)(2,0,2)[7] with drift : Inf
## ARIMA(0,1,2) : 924.7507
## ARIMA(0,1,2) with drift : Inf
## ARIMA(0,1,2)(0,0,1)[7] : 920.0795
## ARIMA(0,1,2)(0,0,1)[7] with drift : Inf
## ARIMA(0,1,2)(0,0,2)[7] : 918.644
## ARIMA(0,1,2)(0,0,2)[7] with drift : Inf
## ARIMA(0,1,2)(1,0,0)[7] : 916.7102
## ARIMA(0,1,2)(1,0,0)[7] with drift : Inf
## ARIMA(0,1,2)(1,0,1)[7] : Inf
## ARIMA(0,1,2)(1,0,1)[7] with drift : Inf
## ARIMA(0,1,2)(1,0,2)[7] : Inf
## ARIMA(0,1,2)(1,0,2)[7] with drift : Inf
## ARIMA(0,1,2)(2,0,0)[7] : 913.8304
## ARIMA(0,1,2)(2,0,0)[7] with drift : Inf
## ARIMA(0,1,2)(2,0,1)[7] : Inf
## ARIMA(0,1,2)(2,0,1)[7] with drift : Inf
## ARIMA(0,1,3) : 926.4351
## ARIMA(0,1,3) with drift : Inf
## ARIMA(0,1,3)(0,0,1)[7] : 920.5319
## ARIMA(0,1,3)(0,0,1)[7] with drift : Inf
## ARIMA(0,1,3)(0,0,2)[7] : 919.5866
## ARIMA(0,1,3)(0,0,2)[7] with drift : Inf
## ARIMA(0,1,3)(1,0,0)[7] : 916.7909
## ARIMA(0,1,3)(1,0,0)[7] with drift : Inf
## ARIMA(0,1,3)(1,0,1)[7] : Inf
## ARIMA(0,1,3)(1,0,1)[7] with drift : Inf
## ARIMA(0,1,3)(2,0,0)[7] : 915.0381
## ARIMA(0,1,3)(2,0,0)[7] with drift : Inf
## ARIMA(0,1,4) : 927.5227
## ARIMA(0,1,4) with drift : Inf
## ARIMA(0,1,4)(0,0,1)[7] : 922.5064
## ARIMA(0,1,4)(0,0,1)[7] with drift : Inf
## ARIMA(0,1,4)(1,0,0)[7] : 918.7093

```

```

## ARIMA(0,1,4)(1,0,0)[7] with drift : Inf
## ARIMA(0,1,5) : 929.2936
## ARIMA(0,1,5) with drift : Inf
## ARIMA(1,1,0) : 944.6207
## ARIMA(1,1,0) with drift : 946.6182
## ARIMA(1,1,0)(0,0,1)[7] : 935.3688
## ARIMA(1,1,0)(0,0,1)[7] with drift : 937.3575
## ARIMA(1,1,0)(0,0,2)[7] : 935.7208
## ARIMA(1,1,0)(0,0,2)[7] with drift : 937.7102
## ARIMA(1,1,0)(1,0,0)[7] : 932.152
## ARIMA(1,1,0)(1,0,0)[7] with drift : 934.1414
## ARIMA(1,1,0)(1,0,1)[7] : Inf
## ARIMA(1,1,0)(1,0,1)[7] with drift : Inf
## ARIMA(1,1,0)(1,0,2)[7] : Inf
## ARIMA(1,1,0)(1,0,2)[7] with drift : Inf
## ARIMA(1,1,0)(2,0,0)[7] : 932.7419
## ARIMA(1,1,0)(2,0,0)[7] with drift : 934.7357
## ARIMA(1,1,0)(2,0,1)[7] : Inf
## ARIMA(1,1,0)(2,0,1)[7] with drift : Inf
## ARIMA(1,1,0)(2,0,2)[7] : Inf
## ARIMA(1,1,0)(2,0,2)[7] with drift : Inf
## ARIMA(1,1,1) : 924.6769
## ARIMA(1,1,1) with drift : Inf
## ARIMA(1,1,1)(0,0,1)[7] : 919.8443
## ARIMA(1,1,1)(0,0,1)[7] with drift : Inf
## ARIMA(1,1,1)(0,0,2)[7] : 918.5174
## ARIMA(1,1,1)(0,0,2)[7] with drift : Inf
## ARIMA(1,1,1)(1,0,0)[7] : 916.4674
## ARIMA(1,1,1)(1,0,0)[7] with drift : Inf
## ARIMA(1,1,1)(1,0,1)[7] : Inf
## ARIMA(1,1,1)(1,0,1)[7] with drift : Inf
## ARIMA(1,1,1)(1,0,2)[7] : Inf
## ARIMA(1,1,1)(1,0,2)[7] with drift : Inf
## ARIMA(1,1,1)(2,0,0)[7] : 913.728
## ARIMA(1,1,1)(2,0,0)[7] with drift : Inf
## ARIMA(1,1,1)(2,0,1)[7] : 915.7907
## ARIMA(1,1,1)(2,0,1)[7] with drift : Inf
## ARIMA(1,1,2) : 926.6681
## ARIMA(1,1,2) with drift : Inf
## ARIMA(1,1,2)(0,0,1)[7] : 921.5654
## ARIMA(1,1,2)(0,0,1)[7] with drift : Inf
## ARIMA(1,1,2)(0,0,2)[7] : 920.3171
## ARIMA(1,1,2)(0,0,2)[7] with drift : Inf
## ARIMA(1,1,2)(1,0,0)[7] : 917.9811
## ARIMA(1,1,2)(1,0,0)[7] with drift : Inf
## ARIMA(1,1,2)(1,0,1)[7] : Inf
## ARIMA(1,1,2)(1,0,1)[7] with drift : Inf
## ARIMA(1,1,2)(2,0,0)[7] : Inf
## ARIMA(1,1,2)(2,0,0)[7] with drift : Inf
## ARIMA(1,1,3) : 928.1161
## ARIMA(1,1,3) with drift : Inf
## ARIMA(1,1,3)(0,0,1)[7] : Inf
## ARIMA(1,1,3)(0,0,1)[7] with drift : Inf
## ARIMA(1,1,3)(1,0,0)[7] : 917.4552

```

```

## ARIMA(1,1,3)(1,0,0)[7] with drift : Inf
## ARIMA(1,1,4) : 929.5955
## ARIMA(1,1,4) with drift : Inf
## ARIMA(2,1,0) : 942.9122
## ARIMA(2,1,0) with drift : 944.9121
## ARIMA(2,1,0)(0,0,1)[7] : 934.9384
## ARIMA(2,1,0)(0,0,1)[7] with drift : 936.9341
## ARIMA(2,1,0)(0,0,2)[7] : 934.606
## ARIMA(2,1,0)(0,0,2)[7] with drift : 936.6031
## ARIMA(2,1,0)(1,0,0)[7] : 931.1616
## ARIMA(2,1,0)(1,0,0)[7] with drift : 933.1577
## ARIMA(2,1,0)(1,0,1)[7] : Inf
## ARIMA(2,1,0)(1,0,1)[7] with drift : Inf
## ARIMA(2,1,0)(1,0,2)[7] : Inf
## ARIMA(2,1,0)(1,0,2)[7] with drift : Inf
## ARIMA(2,1,0)(2,0,0)[7] : 930.6233
## ARIMA(2,1,0)(2,0,0)[7] with drift : 932.6229
## ARIMA(2,1,0)(2,0,1)[7] : Inf
## ARIMA(2,1,0)(2,0,1)[7] with drift : Inf
## ARIMA(2,1,1) : 926.6461
## ARIMA(2,1,1) with drift : Inf
## ARIMA(2,1,1)(0,0,1)[7] : 920.9339
## ARIMA(2,1,1)(0,0,1)[7] with drift : Inf
## ARIMA(2,1,1)(0,0,2)[7] : 919.8506
## ARIMA(2,1,1)(0,0,2)[7] with drift : Inf
## ARIMA(2,1,1)(1,0,0)[7] : 917.0336
## ARIMA(2,1,1)(1,0,0)[7] with drift : Inf
## ARIMA(2,1,1)(1,0,1)[7] : Inf
## ARIMA(2,1,1)(1,0,1)[7] with drift : Inf
## ARIMA(2,1,1)(2,0,0)[7] : 915.1373
## ARIMA(2,1,1)(2,0,0)[7] with drift : Inf
## ARIMA(2,1,2) : 928.3411
## ARIMA(2,1,2) with drift : Inf
## ARIMA(2,1,2)(0,0,1)[7] : Inf
## ARIMA(2,1,2)(0,0,1)[7] with drift : Inf
## ARIMA(2,1,2)(1,0,0)[7] : 916.9772
## ARIMA(2,1,2)(1,0,0)[7] with drift : Inf
## ARIMA(2,1,3) : 927.9445
## ARIMA(2,1,3) with drift : Inf
## ARIMA(3,1,0) : 941.6013
## ARIMA(3,1,0) with drift : 943.5957
## ARIMA(3,1,0)(0,0,1)[7] : 934.2744
## ARIMA(3,1,0)(0,0,1)[7] with drift : 936.2742
## ARIMA(3,1,0)(0,0,2)[7] : 933.4882
## ARIMA(3,1,0)(0,0,2)[7] with drift : 935.4867
## ARIMA(3,1,0)(1,0,0)[7] : 930.248
## ARIMA(3,1,0)(1,0,0)[7] with drift : 932.2476
## ARIMA(3,1,0)(1,0,1)[7] : Inf
## ARIMA(3,1,0)(1,0,1)[7] with drift : Inf
## ARIMA(3,1,0)(2,0,0)[7] : 928.6093
## ARIMA(3,1,0)(2,0,0)[7] with drift : 930.6034
## ARIMA(3,1,1) : 927.1447
## ARIMA(3,1,1) with drift : 926.7629
## ARIMA(3,1,1)(0,0,1)[7] : 922.13

```

```
## ARIMA(3,1,1)(0,0,1)[7] with drift : Inf
## ARIMA(3,1,1)(1,0,0)[7] : 918.6149
## ARIMA(3,1,1)(1,0,0)[7] with drift : Inf
## ARIMA(3,1,2) : 927.2523
## ARIMA(3,1,2) with drift : 927.4237
## ARIMA(4,1,0) : 941.885
## ARIMA(4,1,0) with drift : 943.8584
## ARIMA(4,1,0)(0,0,1)[7] : 935.8141
## ARIMA(4,1,0)(0,0,1)[7] with drift : 937.8106
## ARIMA(4,1,0)(1,0,0)[7] : 931.9403
## ARIMA(4,1,0)(1,0,0)[7] with drift : 933.938
## ARIMA(4,1,1) : 927.7835
## ARIMA(4,1,1) with drift : 927.6555
## ARIMA(5,1,0) : 933.9611
## ARIMA(5,1,0) with drift : 935.8068
##
##
##
## Best model: ARIMA(0,1,1)(2,0,0)[7]
```

```
## Series: lettuce_train
## ARIMA(0,1,1)(2,0,0)[7]
##
## Coefficients:
##          ma1      sar1      sar2
##        -0.9282  0.2638  0.2617
## s.e.    0.0436  0.1074  0.1120
##
## sigma^2 estimated as 1695: log likelihood=-452.19
## AIC=912.38 AICc=912.86 BIC=922.29
```

```
# Best model: ARIMA(0,1,1)(2,0,0)[7] (AIC=912.38)
# Second best: ARIMA(1,1,1)(2,0,0)[7] (AIC=913.72)
# Third best: ARIMA(0,1,2)(2,0,0)[7] (AIC=913.83)
```

Based on the output of `auto.arima()`, a couple of models have similar AICs. Now suppose that we choose the three models with the lowest AICs, namely `ARIMA(0,1,1)(2,0,0)[7]` with `AIC=912.38`, `ARIMA(1,1,1)(2,0,0)[7]` with `AIC=913.72` AND `ARIMA(0,1,2)(2,0,0)[7]` with `AIC=913.83`, as the candidate models that we would like to evaluate further.

```
# three candidate models
lettuce.m1 <- Arima(lettuce_train, order = c(0, 1, 1),
  seasonal = list(order = c(2, 0, 0), period = 7))
lettuce.m2 <- Arima(lettuce_train, order = c(1, 1, 1),
  seasonal = list(order = c(2, 0, 0), period = 7))
lettuce.m3 <- Arima(lettuce_train, order = c(0, 1, 2),
  seasonal = list(order = c(2, 0, 0), period = 7))
```

Now we evaluate the in-sample performance/fit of the model with `accuracy()` function, which summarizes various measures of fitting errors.

A couple of functions are proved to be useful for us to evaluate the in-sample performance/fit of the model. One is `accuracy()` function, which summarizes various measures of fitting errors. In the post-estimation

analysis, we would also like to check out the residual plots, including time series, ACFs and etc, to make sure that there is no warning signal. In particular, residuals shall have a zero mean, constant variance, and distributed symmetrically around mean zero. ACF of any lag greater 0 is expected to be statistically insignificant.

```
# in-sample one-step forecasts model 1  
accuracy(lettuce.m1)
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE  
## Training set 4.103944 40.24048 31.06843 -0.5328407 11.90586 0.8402411  
##              ACF1  
## Training set 0.05766809
```

```
# in-sample one-step forecasts model 2  
accuracy(lettuce.m2)
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE  
## Training set 4.36677 40.13025 30.81424 -0.4249124 11.82139 0.8333666  
##              ACF1  
## Training set -0.02958127
```

```
# in-sample one-step forecasts model 3  
accuracy(lettuce.m3)
```


```
##              ME      RMSE      MAE      MPE      MAPE      MASE  
## Training set 4.317327 40.14598 30.86092 -0.4448752 11.83567 0.8346291  
##              ACF1  
## Training set -0.01581229
```

The second model has the lowest RMSE, even though the first has a lowest AIC score.


Now we proceed with the residual analysis of the three models.

```
# residual analysis model 1  
autoplot(lettuce.m1$residuals, xlab = "Weeks", ylab = "Lettuce") + theme_minimal() +  
ggtitle("New York 1 Store (12631) - Residuals model 1 plot")
```

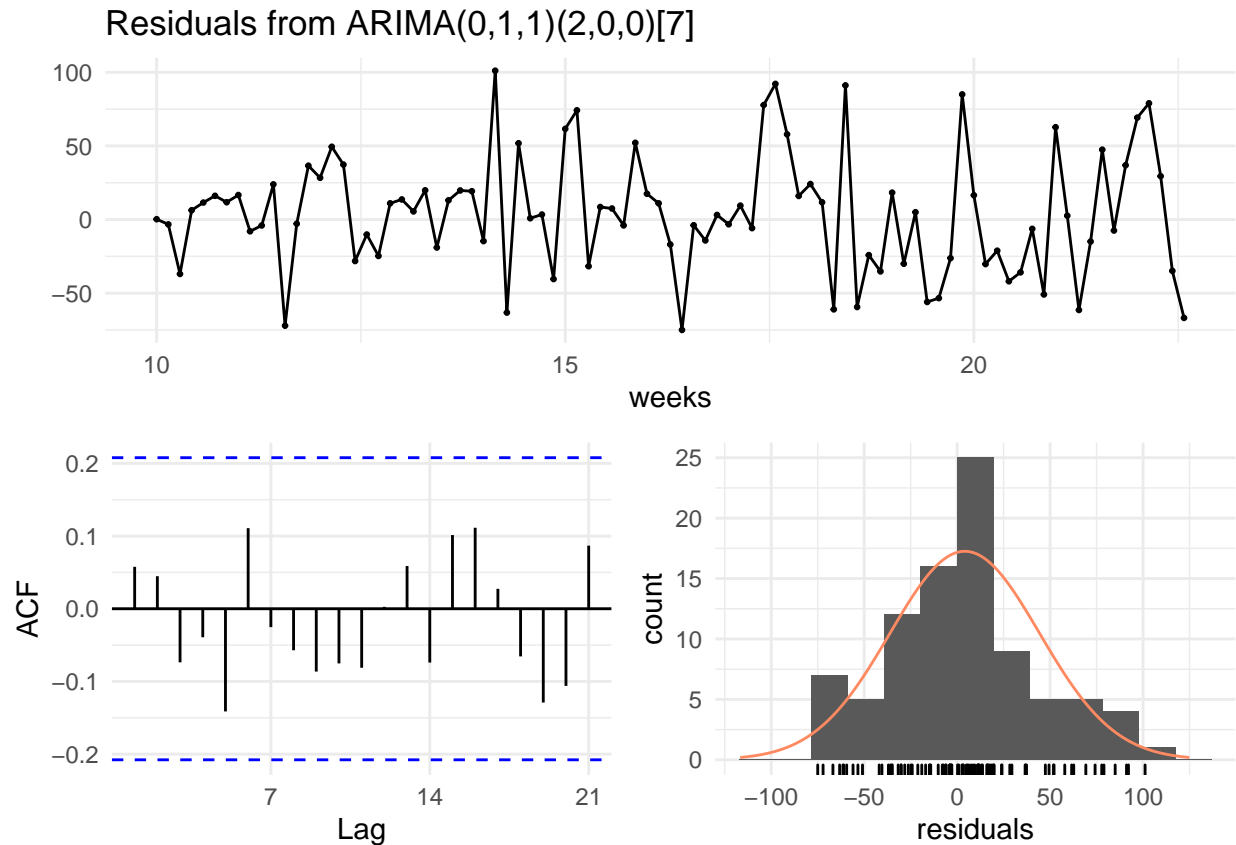
New York 1 Store (12631) – Residuals model 1 plot



```
ggAcf(lettuce.m1$residuals) + theme_minimal() +  
ggtitle("New York 1 Store (12631) - ACF residuals plot model 1")
```




```
checkresiduals(lettuce.m1, xlab = "weeks", theme = theme_minimal())
```

```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,1,1)(2,0,0)[7]
## Q* = 7.6411, df = 11, p-value = 0.745
##
## Model df: 3.   Total lags used: 14
```


```
# residual analysis model 2
autoplot(lettuce.m2$residuals, xlab = "Weeks", ylab = "Lettuce") + theme_minimal() +
ggtitle("New York 1 Store (12631) - Residuals model 2 plot")
```

New York 1 Store (12631) – Residuals model 2 plot




```
ggAcf(lettuce.m2$residuals) + theme_minimal() +  
ggtitle("New York 1 Store (12631) - ACF residuals plot model 2")
```

New York 1 Store (12631) – ACF residuals plot model 2




```
checkresiduals(lettuce.m2, xlab = "weeks", theme = theme_minimal())
```



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(1,1,1)(2,0,0)[7]
## Q* = 7.5634, df = 10, p-value = 0.6714
##
## Model df: 4.   Total lags used: 14
```


```
# residual analysis model 3
autoplot(lettuce.m3$residuals, xlab = "Weeks", ylab = "Lettuce") + theme_minimal() +
ggtitle("New York 1 Store (12631) - Residuals model 3 plot")
```

New York 1 Store (12631) – Residuals model 3 plot




```
ggAcf(lettuce.m3$residuals) + theme_minimal() +  
ggtitle("New York 1 Store (12631) - ACF residuals plot model 3")
```

New York 1 Store (12631) – ACF residuals plot model 3



```
checkresiduals(lettuce.m3, xlab = "weeks", theme = theme_minimal())
```




```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(0,1,2)(2,0,0)[7]
## Q* = 7.5367, df = 10, p-value = 0.674
##
## Model df: 4.    Total lags used: 14
```

Now we continue with the forecasting part for the three candidate models:


```
#Forecasting part model 1
lettuce.f1 <- forecast(lettuce.m1, h = 14)
autoplot(lettuce.f1, xlab = "Weeks", ylab = "Lettuce") + theme_minimal()
```

Forecasts from ARIMA(0,1,1)(2,0,0)[7]




```
#Forecasting part model 2  
lettuce.f2 <- forecast(lettuce.m2, h = 14)  
autoplot(lettuce.f2, xlab = "Weeks", ylab = "Lettuce") + theme_minimal()
```


Forecasts from ARIMA(1,1,1)(2,0,0)[7]



```
#Forecasting part model 3  
lettuce.f3 <- forecast(lettuce.m3, h = 14)  
autoplot(lettuce.f3, xlab = "Weeks", ylab = "Lettuce") + theme_minimal()
```

Forecasts from ARIMA(0,1,2)(2,0,0)[7]



Now we need to test how our models performs for test set. Earlier observations are used for training, and more recent observations are used for testing. Suppose we use the first 89 days of data for training and the last 14 for test. Based on `auto.arima()`, we choose two candidate models with the lowest AICs.

```
### model evaluation
# Apply fitted model to later data
# Accuracy test for candidate model 1
accuracy.m1 <- accuracy(forecast(lettuce.m1, h = 14), lettuce_test)
accuracy.m1
```

```
##
## Training set      ME      RMSE      MAE      MPE      MAPE      MASE
## Test set         -12.488880 41.26388 37.45403 -6.8059927 14.08000 1.0129389
##
## ACF1 Theil's U
## Training set     0.057668088      NA
## Test set         0.009455242 0.6520634
```

```
# Accuracy test for candidate model 2
accuracy.m2 <- accuracy(forecast(lettuce.m2, h = 14), lettuce_test)
accuracy.m2
```

```
##
## Training set      ME      RMSE      MAE      MPE      MAPE      MASE
## Test set         -11.15152 41.18892 36.99544 -6.3085458 13.84815 1.0005363
##
## ACF1 Theil's U
```

```
## Training set -0.029581270      NA
## Test set      0.006889034 0.6535119
```

```
# Accuracy test for candidate model 3
```

```
accuracy.m3 <- accuracy(forecast(lettuce.m3, h = 14), lettuce_test)
accuracy.m3
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  4.317327 40.14598 30.86092 -0.4448752 11.83567 0.8346291
## Test set     -11.480692 41.22657 37.12523 -6.4329773 13.91183 1.0040464
##              ACF1 Theil's U
## Training set -0.015812288      NA
## Test set     0.006505855 0.6536673
```

Thus we pick the second model, since it performs better on the test set.

Now we train the second model on the whole date set as follows:

```
# Training on both train and test set
```

```
lettuce.f.both <- Arima(lettuce, order = c(1, 1, 1),
                        seasonal = list(order = c(2, 0, 0), period = 7))
```

Lastly, we forecast lettuce demand for the next 2 weeks.

```
# Forecast for next 14 days
```

```
lettuce.f.final <- forecast(lettuce.f.both, h = 14)
lettuce.f.final
```

```
##      Point Forecast    Lo 80    Hi 80    Lo 95    Hi 95
## 24.71429      267.8324 214.6772 320.9876 186.5385 349.1263
## 24.85714      252.9936 199.3167 306.6706 170.9018 335.0855
## 25.00000      293.8467 240.0524 347.6410 211.5754 376.1180
## 25.14286      276.2834 222.3930 330.1737 193.8652 358.7015
## 25.28571      253.1750 199.1905 307.1595 170.6129 335.7372
## 25.42857      266.4303 212.3519 320.5086 183.7246 349.1360
## 25.57143      248.8072 194.6352 302.9793 165.9583 331.6562
## 25.71429      277.1748 221.2675 333.0822 191.6720 362.6777
## 25.85714      254.6642 198.5615 310.7669 168.8625 340.4659
## 26.00000      309.5395 253.3018 365.7773 223.5313 395.5478
## 26.14286      302.7129 246.3450 359.0808 216.5057 388.9202
## 26.28571      252.7133 196.2160 309.2106 166.3081 339.1185
## 26.42857      257.0218 200.3953 313.6482 170.4191 343.6244
## 26.57143      252.8329 196.0776 309.5881 166.0332 339.6325
```

We present our forecast through ARIMA(1,1,1)(2,0,0) model for each of the next 14 days.

```
forecast_data <- as.data.frame(lettuce.f.final)
next2weeks <- data.frame(day = seq(1, 14))
final_forecast_NewYork1_arma <- cbind(next2weeks, forecast_data$`Point Forecast`)
final_forecast_NewYork1_arma
```

##	day	forecast_data\$`Point Forecast`
## 1	1	267.8324
## 2	2	252.9936
## 3	3	293.8467
## 4	4	276.2834
## 5	5	253.1750
## 6	6	266.4303
## 7	7	248.8072
## 8	8	277.1748
## 9	9	254.6642
## 10	10	309.5395
## 11	11	302.7129
## 12	12	252.7133
## 13	13	257.0218
## 14	14	252.8329


Holt-Winters

Now we will use another model to forecast lettuce demand. Our goal is to pick the model with the most accurate predictions.

We will forecast the lettuce demand for next two weeks using Holt-Winters model.

For time series analysis, the first step is always to visually inspect the time series. In this regard, the `stl()` function is quite useful. It decomposes the original time series into trend, seasonal factors, and random error terms. The relative importance of different components are indicated by the grey bars in the plots.

```
lettuce_train %>% stl(s.window = "period") %>%
autoplot(xlab = "Weeks", ylab = "Lettuce") + theme_minimal() +
ggtitle("New York 1 Store (12631) - Range bar plot")
```



For this data set, the grey bar of the trend panel is larger than that on the original time series panel, which indicates that the contribution of the trend component to the variation in the original time series is marginal.

Moreover, the grey bar of the seasonal panel is large, even larger than the grey bar of random error term, which indicates that seasonal component contributes to a small proportion of variations in the time series. In other words, it indicates that there is no seasonality in the data.

With `ets()`, initial states and smoothing parameters are jointly estimated by maximizing the likelihood function. We need to specify the model in `ets()` using three letters. The way to approach this is: (1) check out time series plot, and see if there is any trend and seasonality; (2) run `ets()` with `model = "ZZZ"`, and see whether the best model is consistent with your expectation; (3) if they are consistent, it gives us confidence that our model specification is correct; otherwise try to figure out why there is a discrepancy.

We now use `ets` function as previously indicated to find our best model:

```
# using ets
lettuce.ets2 <- ets(lettuce_train, model = "ZZZ")
lettuce.ets2
```

```
## ETS(M,Ad,M)
##
## Call:
## ets(y = lettuce_train, model = "ZZZ")
##
## Smoothing parameters:
##   alpha = 1e-04
##   beta  = 1e-04
##   gamma = 1e-04
```

```
##      phi    = 0.9753
##
##      Initial states:
##      l = 214.8392
##      b = 1.968
##      s = 1.0449 0.9662 0.9145 0.9813 0.853 1.1398
##          1.1003
##
##      sigma: 0.1389
##
##      AIC      AICc      BIC
## 1050.788 1055.641 1083.140
```

Our best model is the ETS(M,Ad,M).

```
# using ets
lettuce.ets <- ets(lettuce_train, model = "MAM", damped = TRUE, ic = 'aic')
lettuce.ets
```

```
## ETS(M,Ad,M)
##
## Call:
## ets(y = lettuce_train, model = "MAM", damped = TRUE, ic = "aic")
##
##      Smoothing parameters:
##      alpha = 1e-04
##      beta  = 1e-04
##      gamma = 1e-04
##      phi   = 0.9753
##
##      Initial states:
##      l = 214.8392
##      b = 1.968
##      s = 1.0449 0.9662 0.9145 0.9813 0.853 1.1398
##          1.1003
##
##      sigma: 0.1389
##
##      AIC      AICc      BIC
## 1050.788 1055.641 1083.140
```

After estimation, we can use `accuracy()` function to determine in-sample fit and `forecast()` function to generate forecast.

Similarly with ARIMA model, we use AIC to determine our best model in terms of best in-sample performance.

```
# in-sample one-step forecast
accuracy(lettuce.ets)
```

```
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -1.077097 34.11242 26.09866 -2.108286 10.08561 0.7058346
##              ACF1
## Training set 0.06152982
```

We present the in-sample forecast part for the ets model as follows:

```
# best model
lettuce.ets.f <- forecast(lettuce.ets, h = 14)
lettuce.ets.f
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
## 22.71429	274.6589	225.7799	323.5380	199.9049	349.4130
## 22.85714	297.2379	244.3406	350.1352	216.3385	378.1373
## 23.00000	313.2411	257.4958	368.9863	227.9861	398.4961
## 23.14286	324.6814	266.9002	382.4627	236.3127	413.0502
## 23.28571	243.1538	199.8815	286.4261	176.9745	309.3331
## 23.42857	279.8874	230.0778	329.6969	203.7102	356.0645
## 23.57143	261.0075	214.5579	307.4571	189.9689	332.0461
## 23.71429	275.9276	226.8227	325.0325	200.8282	351.0270
## 23.85714	298.5760	245.4406	351.7115	217.3123	379.8397
## 24.00000	314.6155	258.6256	370.6055	228.9863	400.2447
## 24.14286	326.0700	268.0415	384.0984	237.3231	414.8168
## 24.28571	244.1674	200.7145	287.6202	177.7120	310.6227
## 24.42857	281.0245	231.0124	331.0366	204.5377	357.5114
## 24.57143	262.0411	215.4074	308.6749	190.7210	333.3613

After the forecast, we continue with the out of sample accuracy of our best model.

```
# Out of sample accuracy
# best model
accuracy.ets <- accuracy(lettuce.ets.f, lettuce_test)
accuracy.ets
```

	ME	RMSE	MAE	MPE	MAPE	MASE
## Training set	-1.077097	34.11242	26.09866	-2.108286	10.08561	0.7058346
## Test set	-19.235015	44.30622	36.85583	-9.276175	14.25979	0.9967607

	ACF1	Theil's U
## Training set	0.061529822	NA
## Test set	-0.001611319	0.7245148

We now train our best model - ETS(M,Ad,M) on the whole data set as indicated below:

```
# final model
lettuce.ets <- ets(lettuce, model = "MAM", damped = TRUE, ic = 'aic')
lettuce.ets
```

```
## ETS(M,Ad,M)
##
## Call:
## ets(y = lettuce, model = "MAM", damped = TRUE, ic = "aic")
##
## Smoothing parameters:
##   alpha = 1e-04
##   beta  = 1e-04
##   gamma = 0.002
##   phi   = 0.9745
```

```
##
## Initial states:
## l = 213.5467
## b = 1.9142
## s = 1.028 0.9728 0.9128 0.9858 0.8681 1.1111
##      1.1214
##
## sigma: 0.1408
##
##      AIC      AICc      BIC
## 1232.106 1236.196 1266.357
```

We now present the out-of-sample forecast for the next 14 days (2 weeks) as seen below:

```
lettuce.ets.f <- forecast(lettuce.ets, h = 14)
lettuce.ets.f
```

##		Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
##	24.71429	274.2683	224.7805	323.7561	198.5832	349.9534
##	24.85714	289.9449	237.6285	342.2614	209.9338	369.9560
##	25.00000	316.5410	259.4257	373.6564	229.1907	403.8914
##	25.14286	313.7897	257.1708	370.4086	227.1986	400.3808
##	25.28571	245.1401	200.9080	289.3721	177.4930	312.7872
##	25.42857	278.4584	228.2145	328.7023	201.6170	355.2998
##	25.57143	257.8716	211.3423	304.4009	186.7112	329.0320
##	25.71429	275.0112	225.3891	324.6332	199.1207	350.9016
##	25.85714	290.7099	238.2552	343.1646	210.4873	370.9325
##	26.00000	317.3545	260.0921	374.6169	229.7792	404.9298
##	26.14286	314.5752	257.8143	371.3362	227.7668	401.3836
##	26.28571	245.7378	201.3976	290.0780	177.9254	313.5503
##	26.42857	279.1198	228.7563	329.4834	202.0954	356.1442
##	26.57143	258.4683	211.8310	305.1055	187.1427	329.7938

We present our forecast for each of the next 14 days.

```
forecast_data <- as.data.frame(lettuce.ets.f)
next2weeks <- data.frame(day = seq(1, 14))
final_forecast_NewYork1_ets <- cbind(next2weeks, forecast_data$`Point Forecast`)
final_forecast_NewYork1_ets
```


##	day	forecast_data\$`Point Forecast`
##	1	274.2683
##	2	289.9449
##	3	316.5410
##	4	313.7897
##	5	245.1401
##	6	278.4584
##	7	257.8716
##	8	275.0112
##	9	290.7099
##	10	317.3545
##	11	314.5752
##	12	245.7378
##	13	279.1198
##	14	258.4683

Comparison

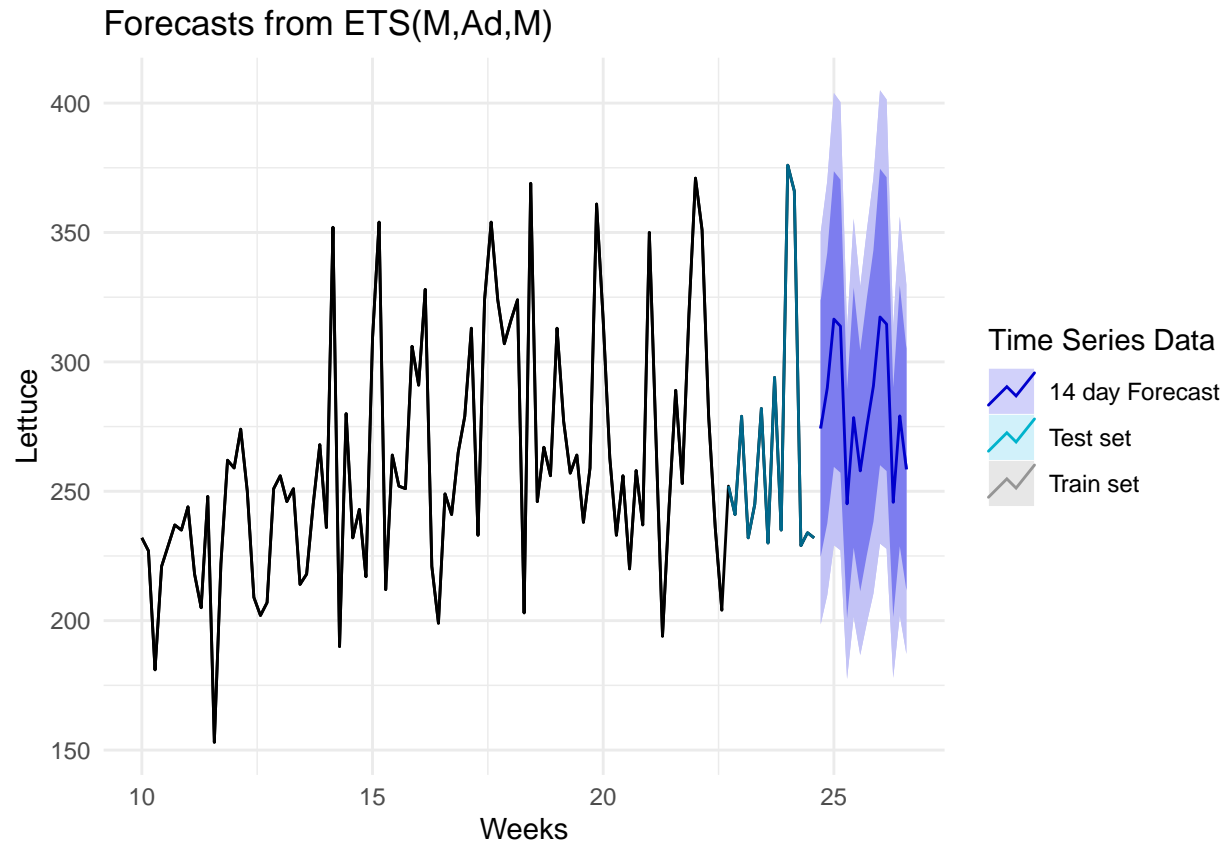
Now we will compare the two best models for New York 1 Store (12631).

We plot time series data for train and test set and also the forecasts from our two models as indicated below:

```
colours <- c("blue", "deepskyblue4", "black")
autoplot(lettuce.f.final, xlab = "Weeks", ylab = "Lettuce") +
  autolayer(lettuce_train, series = "Train set") +
  autolayer(lettuce_test, series = "Test set") +
  autolayer(lettuce.f.final, series = "14 day Forecast") +
  guides(colour = guide_legend(title = "Time Series Data")) +
  scale_colour_manual(values = colours) + theme_minimal()
```



```
autoplot(lettuce.ets.f, xlab = "Weeks", ylab = "Lettuce") +
  autolayer(lettuce_train, series = "Train set") +
  autolayer(lettuce_test, series = "Test set") +
  autolayer(lettuce.ets.f, series = "14 day Forecast") +
  guides(colour = guide_legend(title = "Time Series Data")) +
  scale_colour_manual(values = colours) + theme_minimal()
```



In order to decide which of the two models $ARIMA(1,1,1)(2,0,0)$ or $ETS(M,Ad,M)$ to choose, we will check their RMSE in the test set.

```
# best ets model
# ETS(M,Ad,M)
accuracy.ets
```

```
##                ME      RMSE      MAE      MPE      MAPE      MASE
## Training set  -1.077097 34.11242 26.09866 -2.108286 10.08561 0.7058346
## Test set      -19.235015 44.30622 36.85583 -9.276175 14.25979 0.9967607
##                ACF1 Theil's U
## Training set   0.061529822      NA
## Test set       -0.001611319 0.7245148
```

```
# best arima model
# ARIMA(1,1,1)(2,0,0)
accuracy.m2
```

```
##                ME      RMSE      MAE      MPE      MAPE      MASE
## Training set   4.36677 40.13025 30.81424 -0.4249124 11.82139 0.8333666
## Test set       -11.15152 41.18892 36.99544 -6.3085458 13.84815 1.0005363
##                ACF1 Theil's U
## Training set  -0.029581270      NA
## Test set       0.006889034 0.6535119
```

We can observe that ARIMA(1,1,1)(2,0,0) has a better (lower) RMSE (41.18892 vs 44.30622) respectively. Therefore, we choose the ARIMA(1,1,1)(2,0,0) for New York 1 (12631) store. Hence, our forecast for lettuce demand of next 2 weeks for that store is the following:

```
final_forecast_NewYork1_arima
```

```
##      day forecast_data$`Point Forecast`
## 1      1                267.8324
## 2      2                252.9936
## 3      3                293.8467
## 4      4                276.2834
## 5      5                253.1750
## 6      6                266.4303
## 7      7                248.8072
## 8      8                277.1748
## 9      9                254.6642
## 10    10                309.5395
## 11    11                302.7129
## 12    12                252.7133
## 13    13                257.0218
## 14    14                252.8329
```